

Solution for Merge Sort Comps

Time Complexity (Part 2) and Proof of Correctness (Part 3)

October 20, 2025

Part 2: Time Complexity Analysis

Step-by-Step Cost Analysis

Let n be the number of rows (people).

- **Data Load** Reading the CSV and constructing n objects is $O(n)$. Computing `new_decile` = $0.6P + 0.5J + 0.1C$ per row is $O(1)$, for $O(n)$ total.
- **Recursive Splitting.** Each call partitions its subarray in $O(1)$ time; there are $2n - 1$ calls overall across the recursion tree, so the aggregate splitting overhead is $O(n)$.
- **Merging.** If two sublists have combined size k , a single `merge` performs $k - 1$ comparisons and k moves (both $O(k)$). Summed across one level of the recursion tree, this is $O(n)$ work.

Overall Runtime Discussion

Let $T(n)$ be the worst-case running time of merge sort on n elements. Then

$$T(n) = \begin{cases} \Theta(1), & n \leq 1, \\ 2T\left(\frac{n}{2}\right) + \Theta(n), & n > 1. \end{cases}$$

By the Master Theorem ($a = 2$, $b = 2$, $f(n) = \Theta(n)$, and $n^{\log_b a} = n$),

$$T(n) = \Theta(n \log n).$$

Part 3: Proof of Correctness

Lemma 1 (Merge Correctness)

If L and R are sorted by `comes_before`, then `merge(L,R)` returns a list M such that

- (i) M contains exactly the elements of $L \uplus R$ (multiset union), and
- (ii) M is sorted by `comes_before`.

Proof.

Property i

The merge algorithm maintains two pointers (indices) i and j for lists L and R respectively, both initially 0. At each iteration, the algorithm compares $L[i]$ and $R[j]$ using `comes_before`, appends the smaller element to output list M , and advances only the corresponding pointer. This process continues until one list is exhausted, at which point all remaining elements from the other list are appended to M .

Since each iteration removes exactly one element from either L or R and appends it to M , and the algorithm terminates only when both lists are fully consumed (when $i = |L|$ and $j = |R|$), every element from both input lists appears exactly once in M . No elements are lost and none are duplicated. Therefore, M contains precisely the multiset union $L \uplus R$.

Property ii

We prove by induction on t , the number of elements appended to M , that M maintains the sorted property as a prefix invariant.

Base case ($t = 0$): Initially, M is empty, which is trivially sorted.

Inductive step: Assume after t steps, $M[0..t-1]$ is sorted. At step $t+1$, we append either $L[i]$ or $R[j]$ to M , specifically the element e for which `comes_before`(e, \cdot) returns true when compared with the other list's current head.

Since both L and R are sorted, we have:

- All elements $L[i'], i' \geq i$ in L satisfy $L[i] \leq L[i']$
- All elements $R[j'], j' \geq j$ in R satisfy $R[j] \leq R[j']$

The element e appended at step $t+1$ is the minimum among $\{L[i], R[j]\}$ (the current heads). Since all remaining elements in both lists are \geq their respective heads, e is the minimum among all unprocessed elements.

By the inductive hypothesis, $M[0..t-1]$ is sorted. The last element $M[t-1]$ was chosen in a previous step as the minimum among available elements at that time. Since e is now chosen as the minimum among currently available elements, and all elements previously chosen were from a set that has only grown smaller, we have $M[t-1] \leq e$. Therefore, $M[0..t]$ remains sorted.

When one list is exhausted, the remainder of the other list is appended. Since that list is already sorted and all its remaining elements are \geq the last element added to M by the same reasoning, sortedness is preserved.

In conclusion, by induction, M is sorted throughout the merge process, and thus the final output is sorted. \square

Lemma 2 (Base Case Correctness)

If $|data| \leq 1$, then `merge_sort(data, key)` returns a sorted list containing exactly the input elements.

Proof. A list of size 0 or 1 is vacuously sorted and returned unchanged. \square

Lemma 3 (Merge Sort Correctness)

For any finite list A of `Person` records and any valid key, `merge_sort(A, key)` returns a list that (1) is a permutation of A and (2) is sorted in nonincreasing order by that key, breaking ties by name ascending.

Proof. By strong induction on $n = |A|$.

Base: For $n \leq 1$, the claim holds by Lemma 2.

Inductive step: Assume the claim holds for all lists of size $< n$ (with $n \geq 2$). Split A into two halves L and R with $|L|, |R| < n$. By the inductive hypothesis, `merge_sort(L, key)` and `merge_sort(R, key)` return sorted lists that are permutations of L and R , respectively. Applying Lemma 1 to these two sorted outputs yields a merged list M that is sorted and contains exactly the multiset union of L and R , which is precisely A . Thus M is a sorted permutation of A .

Therefore, by induction, the theorem holds for all n . \square

Termination

The algorithm terminates because:

1. The base case ($n \leq 1$) performs no recursive calls
2. For $n \geq 2$, each recursive call operates on a strictly smaller list ($\lfloor n/2 \rfloor < n$ and $\lceil n/2 \rceil < n$)
3. The merge operation processes elements linearly and always terminates in $O(|L| + |R|)$ steps